

## Task 1

### Problem 1

```
yanbo@yanbo-ThinkPad-Edge:~/443/lab4$ gcc -Wall -o sniffex sniffex.c -lpcap
sniffex.c: In function 'got_packet':
sniffex.c:486:10: warning: pointer targets in assignment differ in signedness [-Wpointer-sign]
    payload = (u_char *) (packet + SIZE_ETHERNET + size_ip + size_tcp);
               ^
sniffex.c:497:17: warning: pointer targets in passing argument 1 of 'print_payload' differ in signedness [-Wpointer-sign]
    print_payload(payload, size_payload);
               ^
sniffex.c:373:1: note: expected 'const u_char * {aka const unsigned char *}' but argument is of type 'const char *'
    print_payload(const u_char *payload, int len)
               ^
sniffex.c:424:31: warning: variable 'ethernet' set but not used [-Wunused-but-set-variable]
    const struct sniff_ethernet *ethernet; /* The ethernet header [1] */
               ^
yanbo@yanbo-ThinkPad-Edge:~/443/lab4$ sudo ./sniffex
sniffex - Sniffer example using libpcap
Copyright (c) 2005 The Tcpdump Group
THERE IS ABSOLUTELY NO WARRANTY FOR THIS PROGRAM.

Device: wlp3s0
Number of packets: 10
Filter expression: ip

Packet number 1:
    From: 172.217.8.166
          To: 104.38.187.117
    Protocol: TCP
    Src port: 443
    Dst port: 59778
    Payload (55 bytes):
00000  17 03 03 00 32 1e 34 93  d3 44 ab fd 96 91 9e 6f  ....2.4..D....0
00016  1c 73 ec ff eb 3f 7c d1  e9 a2 b0 57 0b 19 66 73  .s...?|....W..fs
00032  ca dd db d9 ba 64 be ed  39 68 94 95 cc 28 81 5c  .....d..9h...(.\
00048  c2 ab 38 3a cf f9 3b                                ..8:...;

Packet number 2:
    From: 172.217.8.166
          To: 104.38.187.117
    Protocol: TCP
    Src port: 443
```

```
To: 151.101.129.69
Protocol: TCP
Src port: 36778
Dst port: 443

Packet number 6:
    From: 104.38.187.117
        To: 151.101.129.69
    Protocol: TCP
    Src port: 36780
    Dst port: 443

Packet number 7:
    From: 104.38.187.117
        To: 151.101.129.69
    Protocol: TCP
    Src port: 36782
    Dst port: 443

Packet number 8:
    From: 172.217.8.166
        To: 104.38.187.117
    Protocol: TCP
    Src port: 443
    Dst port: 59778

Packet number 9:
    From: 151.101.193.69
        To: 104.38.187.117
    Protocol: TCP
    Src port: 443
    Dst port: 45078

Packet number 10:
    From: 151.101.129.69
        To: 104.38.187.117
    Protocol: TCP
    Src port: 443
    Dst port: 36778

Capture complete.
```

The following calls are functions of sniffex.c:

Pcap\_lookupdev is to find a capture device to sniff.

Pcap\_lookupnet returns the network number and mask of the capture device.

Pcap\_open\_live starts sniffing on the capture device.

Pcap\_datalink returns the device we are capturing.

Pcap\_compile compiles a filter expression stored in a regular string to set the filter.

Pcap\_setfilter sets the compiled filter. We can sniff a packet (pcap\_next) or continuously sniff (pcap\_loop). We can either sniff one packet at a time (pcap\_next) or continuously sniff (pcap\_loop). We will continue with pcap\_loop: Sets callback function for new (filtered!) packets

Pcap\_freecode frees the memory generated by pcap\_compile.

Pcap\_close closes the sniffer session.

## Problem 2

```
yanbo@yanbo-ThinkPad-Edge:~/443/Lab4$ ./sniffex
sniffex - Sniffer example using libpcap
Copyright (c) 2005 The Tcpdump Group
THERE IS ABSOLUTELY NO WARRANTY FOR THIS PROGRAM.

Device: wlp3s0
[SMPlayer] :kets: 10
Filter expression: ip
Couldn't open device wlp3s0: wlp3s0: You don't have permission to capture on that device (socket: Operation not permitted)
```

We need root for sniffex to run because sniffex will need to access a network device and non root user cannot do that.

The code causing failure is:

```
dev = pcap_lookupdev(errbuf);
if (dev == NULL) {
    fprintf(stderr, "Couldn't find default device: %s\n", errbuf);
    exit(EXIT_FAILURE);
}
```

## Problem 3

Hybrid mode allows network sniffers to transfer all traffic from network controllers, not just traffic to the network controller. On pcap\_open\_live, the capture device determines the third parameter (a Boolean int) in the mixed mode in line 551. The following code shows the differences.

```
/* promisc mode on */
handle = pcap_open_live(dev, SNAP_LEN, 1, 1000, errbuf);

/* promisc mode off */
handle = pcap_open_live(dev, SNAP_LEN, 0, 1000, errbuf);
```

By default, sniffex scoffs at hybrid models. To see the difference in IP addresses, we switched the VM's network adapter to "bridge" and then ran sniffex. First, we get a pop-up window from VMWare Fusion to ask if the virtual machine can monitor all traffic which we allow it.



william@ubuntu: ~/lab4

```
Wpointer-sign]
    payload = (u_char *)(packet + SIZE_ETHERNET + size_ip + size_tcp);

sniffex.c
ad' differ
    print_
sniffex.c
argument
print_pa
^
sniffex.c
t-variabl
const s
william@u
sniffex -
Copyright (c) 2005 The Tcpdump Group
THERE IS ABSOLUTELY NO WARRANTY FOR THIS PROGRAM.

Device: ens33
Number of packets: 10
Filter expression: ip
```

A virtual machine is attempting to monitor all network traffic, which requires administrator access.  
Enter your password to allow this.

User Name: William.Sun  
Password:

Cancel OK

As we allow it, sniffex picks up more packets from the host machine. It is shown below:

```
william@ubuntu:~/lab4$ ifconfig
ens33      Link encap:Ethernet HWaddr 00:0c:29:f9:cc:40
            inet addr:104.39.142.37  Bcast:104.39.255.255  Mask:255.255.0.0
            inet6 addr: fe80::1a96:ed10:a66c:8f5/64 Scope:Link
                  UP BROADCAST RUNNING MULTICAST  MTU:1500  Metric:1
                  RX packets:287610 errors:0 dropped:0 overruns:0 frame:0
                  TX packets:120603 errors:0 dropped:0 overruns:0 carrier:0
                  collisions:0 txqueuelen:1000
                  RX bytes:391601643 (391.6 MB)  TX bytes:6591329 (6.5 MB)
                  Interrupt:19 Base address:0x2000

lo         Link encap:Local Loopback
            inet addr:127.0.0.1  Mask:255.0.0.0
            inet6 addr: ::1/128 Scope:Host
                  UP LOOPBACK RUNNING  MTU:65536  Metric:1
                  RX packets:431 errors:0 dropped:0 overruns:0 frame:0
                  TX packets:431 errors:0 dropped:0 overruns:0 carrier:0
                  collisions:0 txqueuelen:1000
                  RX bytes:34465 (34.4 KB)  TX bytes:34465 (34.4 KB)

william@ubuntu:~/lab4$ sudo ./sniffex
sniffex - Sniffer example using libpcap
Copyright (c) 2005 The Tcpdump Group
THERE IS ABSOLUTELY NO WARRANTY FOR THIS PROGRAM.

Device: ens33
Number of packets: 10
Filter expression: ip

Packet number 1:
    From: 104.39.76.237
          To: 74.125.124.189
    Protocol: TCP
    Src port: 61442
    Dst port: 443

Packet number 2:
    From: 104.39.76.237
          To: 74.125.124.189
    Protocol: TCP
    Src port: 61442
    Dst port: 443

Packet number 3:
    From: 104.39.76.237
          To: 74.125.124.189
    Protocol: TCP
    Src port: 61442
```

```

[WilliamSundeMacBook-Pro:~ william.sun$ ifconfig
lo0: flags=8049<UP,LOOPBACK,RUNNING,MULTICAST> mtu 16384
    options=1203<RXCSUM,TXCSUM,TXSTATUS,SW_TIMESTAMP>
    inet 127.0.0.1 netmask 0xffff000000
    inet6 ::1 prefixlen 128
    inet6 fe80::1%lo0 prefixlen 64 scopeid 0x1
        nd6 options=201<PERFORMNUD,DAD>
gif0: flags=8010<POINTOPOINT,MULTICAST> mtu 1280
stf0: flags=0<> mtu 1280
XHC1: flags=0<> mtu 0
XHC20: flags=0<> mtu 0
XHC0: flags=0<> mtu 0
en0: flags=8863<UP,BROADCAST,SMART,RUNNING,SIMPLEX,MULTICAST> mtu 1500
    ether 78:4f:43:74:eb:16
    inet6 fe80::c95:828f:becd:591e%en0 prefixlen 64 secured scopeid 0x8
        inet 104.39.76.237 netmask 0xffff0000 broadcast 104.39.255.255
        nd6 options=201<PERFORMNUD,DAD>
    media: autoselect
    status: active
p2p0: flags=8843<UP,BROADCAST,RUNNING,SIMPLEX,MULTICAST> mtu 2304
    ether 0a:4f:43:74:eb:16
    media: autoselect
    status: inactive
awdl0: flags=8943<UP,BROADCAST,RUNNING,PROMISC,SIMPLEX,MULTICAST> mtu 1484
    ether 52:ce:a0:25:97:3b
    inet6 fe80::50ce:a0ff:fe25:973b%awdl0 prefixlen 64 scopeid 0xa
        nd6 options=201<PERFORMNUD,DAD>
    media: autoselect
    status: active
en3: flags=8963<UP,BROADCAST,SMART,RUNNING,PROMISC,SIMPLEX,MULTICAST> mtu 1500
    options=60<TS04,TS06>
    ether 12:00:ed:22:31:01
    media: autoselect <full-duplex>
    status: inactive
en1: flags=8963<UP,BROADCAST,SMART,RUNNING,PROMISC,SIMPLEX,MULTICAST> mtu 1500
    options=60<TS04,TS06>
    ether 12:00:ed:22:31:00
    media: autoselect <full-duplex>
    status: inactive
en4: flags=8963<UP,BROADCAST,SMART,RUNNING,PROMISC,SIMPLEX,MULTICAST> mtu 1500
    options=60<TS04,TS06>
    ether 12:00:ed:22:31:05
    media: autoselect <full-duplex>
    status: inactive
en2: flags=8963<UP,BROADCAST,SMART,RUNNING,PROMISC,SIMPLEX,MULTICAST> mtu 1500
    options=60<TS04,TS06>
```

Then we turn off promiscuous mode from 1 to 0 in pcap\_open\_live. And we run ping 8.8.8.8 and create some network traffic and the results are:

```
william@ubuntu:~/Lab4$ sudo ./sniffex
sniffex - Sniffer example using libpcap
Copyright (c) 2005 The Tcpdump Group
THERE IS ABSOLUTELY NO WARRANTY FOR THIS PROGRAM.

Device: ens33
Number of packets: 10
Filter expression: ip

Packet number 1:
    From: 104.39.142.37
        To: 8.8.8.8
    Protocol: ICMP

Packet number 2:
    From: 8.8.8.8
        To: 104.39.142.37
    Protocol: ICMP

Packet number 3:
    From: 104.39.142.37
        To: 8.8.8.8
    Protocol: ICMP

Packet number 4:
    From: 8.8.8.8
        To: 104.39.142.37
    Protocol: ICMP

Packet number 5:
    From: 104.39.142.37
        To: 8.8.8.8
    Protocol: ICMP

Packet number 6:
    From: 8.8.8.8
        To: 104.39.142.37
    Protocol: ICMP

Packet number 7:
    From: 104.39.142.37
        To: 8.8.8.8
    Protocol: ICMP

Packet number 8:
    From: 8.8.8.8
        To: 104.39.142.37
    Protocol: ICMP
```

### Task 1.b(1)

We filter only the ICMP packets by modifying the `filter_exp` string by using  
`char filter_exp[] = "icmp and (src host 104.39.142.37 and dst host 8.8.8.8) or (src host 8.8.8.8 and dst host 104.39.142.37);`

Then we run the sniffex and ping 8.8.8.8 and got the following results:

william@ubuntu:~/lab4\$ sudo ./sniffex  
 sniffex - Sniffer example using libpcap  
 Copyright (c) 2005 The Tcpdump Group  
 THERE IS ABSOLUTELY NO WARRANTY FOR THIS PROGRAM.  
 Device: ens33  
 Number of packets: 10  
 Filter expression: icmp and (src host 104.39.142.37 and dst host 8.8.8.8) or (src host 8.8.8.8 and dst host 104.39.142.37)  
 Packet number 1:  
 From: 104.39.142.37  
 To: 8.8.8.8  
 Protocol: ICMP  
 Packet number 2:  
 From: 8.8.8.8  
 To: 104.39.142.37  
 Protocol: ICMP  
 Packet number 3:  
 From: 104.39.142.37  
 To: 8.8.8.8  
 Protocol: ICMP  
 Packet number 4:  
 From: 8.8.8.8  
 To: 104.39.142.37  
 Protocol: ICMP  
 Packet number 5:  
 From: 104.39.142.37  
 To: 8.8.8.8  
 Protocol: ICMP  
 Packet number 6:  
 From: 8.8.8.8  
 To: 104.39.142.37  
 Protocol: ICMP  
 Packet number 7:  
 From: 104.39.142.37  
 To: 8.8.8.8  
 Protocol: ICMP  
 Packet number 8:  
 From: 8.8.8.8

william@ubuntu:~/lab4\$ ping 8.8.8.8  
 PING 8.8.8.8 (8.8.8.8) 56(84) bytes of data.  
 64 bytes from 8.8.8.8: icmp\_seq=1 ttl=59 time=18.8 ms  
 64 bytes from 8.8.8.8: icmp\_seq=2 ttl=58 time=18.7 ms  
 64 bytes from 8.8.8.8: icmp\_seq=3 ttl=58 time=18.5 ms  
 64 bytes from 8.8.8.8: icmp\_seq=4 ttl=58 time=18.8 ms  
 64 bytes from 8.8.8.8: icmp\_seq=5 ttl=58 time=17.8 ms  
 64 bytes from 8.8.8.8: icmp\_seq=6 ttl=58 time=20.1 ms  
 64 bytes from 8.8.8.8: icmp\_seq=7 ttl=58 time=18.5 ms  
 64 bytes from 8.8.8.8: icmp\_seq=8 ttl=58 time=18.7 ms  
 64 bytes from 8.8.8.8: icmp\_seq=9 ttl=58 time=18.7 ms  
 64 bytes from 8.8.8.8: icmp\_seq=10 ttl=58 time=18.9 ms  
 64 bytes from 8.8.8.8: icmp\_seq=11 ttl=58 time=18.6 ms  
 64 bytes from 8.8.8.8: icmp\_seq=12 ttl=58 time=18.9 ms  
 64 bytes from 8.8.8.8: icmp\_seq=13 ttl=58 time=18.2 ms  
 ^C  
 --- 8.8.8.8 ping statistics ---  
 13 packets transmitted, 13 received, 0% packet loss, time 12024ms  
 rtt min/avg/max/mdev = 17.829/18.753/20.178/0.515 ms  
 william@ubuntu:~/lab4\$

Save

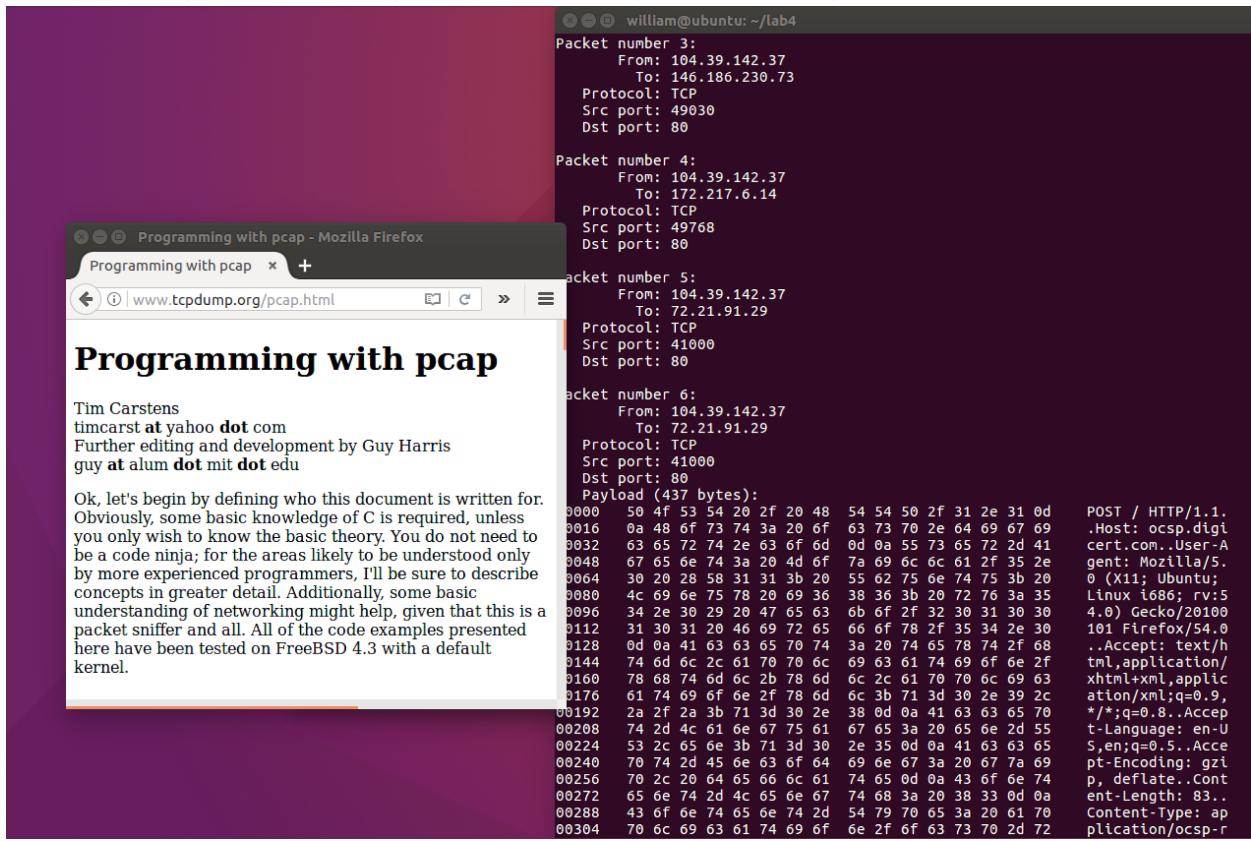
Tab Width: 8 ▾ Ln 522, Col 10 ▾ INS

### Task 1.b (2)

We change the filter\_exp variable and only sniff TCP packets with a destination port from 10 to 100.

```
char filter_exp[] = "tcp dst portrange 10-100";
```

Then we navigate to a webpage in browser and got the following:

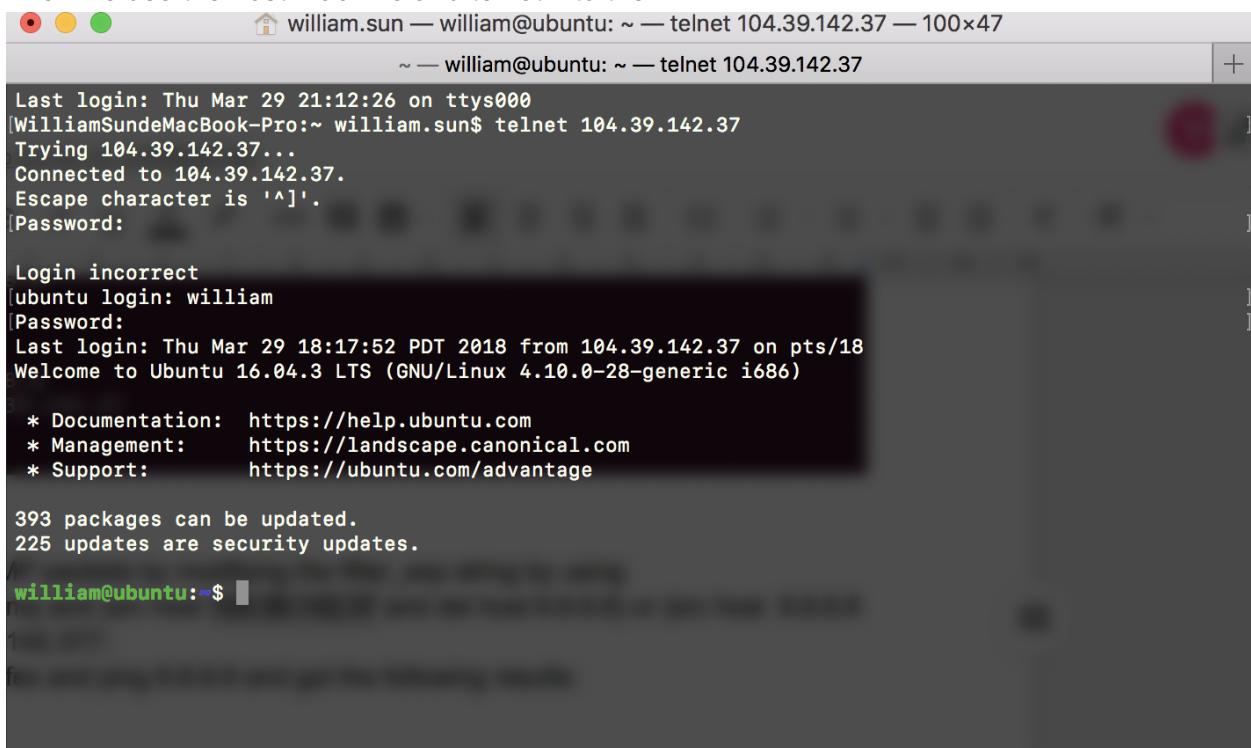


### Task 1.c

We have sniffed telnet passwords, we can just look for tcp packets on port 23.

```
char filter_exp[] = "tcp port 23";
```

Then we use the host machine and telnet into the VM.



```
william.sun — william@ubuntu: ~ — telnet 104.39.142.37 — 100x47
~ — william@ubuntu: ~ — telnet 104.39.142.37

Last login: Thu Mar 29 21:12:26 on ttys000
[WilliamSundeMacBook-Pro:~ william.sun$ telnet 104.39.142.37
Trying 104.39.142.37...
Connected to 104.39.142.37.
Escape character is '^]'.
>Password:

Login incorrect
ubuntu login: william
>Password:
Last login: Thu Mar 29 18:17:52 PDT 2018 from 104.39.142.37 on pts/18
Welcome to Ubuntu 16.04.3 LTS (GNU/Linux 4.10.0-28-generic i686)

 * Documentation:  https://help.ubuntu.com
 * Management:     https://landscape.canonical.com
 * Support:        https://ubuntu.com/advantage

393 packages can be updated.
225 updates are security updates.

william@ubuntu:~$
```

```
Device: ens33
Number of packets: 10
Filter expression: tcp port 23

Packet number 1:
  From: 104.39.76.237
    To: 104.39.142.37
  Protocol: TCP
  Src port: 49456
  Dst port: 23

Packet number 2:
  From: 104.39.142.37
    To: 104.39.76.237
  Protocol: TCP
  Src port: 23
  Dst port: 49456

Packet number 3:
  From: 104.39.76.237
    To: 104.39.142.37
  Protocol: TCP
  Src port: 49456
  Dst port: 23

Packet number 4:
  From: 104.39.76.237
    To: 104.39.142.37
  Protocol: TCP
  Src port: 49456
  Dst port: 23
  Payload (27 bytes):
00000  ff fb 25 ff fd 03 ff fb 18 ff fb 1f ff fb 20 ff  ..%..... .
00016  fb 21 ff fb 22 ff fb 27 ff fd 05  ...."..."...

Packet number 5:
  From: 104.39.142.37
    To: 104.39.76.237
  Protocol: TCP
  Src port: 23
  Dst port: 49456

Packet number 6:
  From: 104.39.142.37
    To: 104.39.76.237
  Protocol: TCP
  Src port: 23
  Dst port: 49456
  Payload (12 bytes):
00000  ff fd 18 ff fd 20 ff fd 23 ff fd 27  ..... .#..'
```

```

Packet number 6:
    From: 104.39.142.37
        To: 104.39.76.237
    Protocol: TCP
    Src port: 23
    Dst port: 49456
    Payload (12 bytes):
00000  ff fd 18 ff fd 20 ff fd 23 ff fd 27 ..... .#..'

Packet number 7:
    From: 104.39.76.237
        To: 104.39.142.37
    Protocol: TCP
    Src port: 49456
    Dst port: 23

Packet number 8:
    From: 104.39.76.237
        To: 104.39.142.37
    Protocol: TCP
    Src port: 49456
    Dst port: 23
    Payload (3 bytes):
00000  ff fc 23 ..#

Packet number 9:
    From: 104.39.142.37
        To: 104.39.76.237
    Protocol: TCP
    Src port: 23
    Dst port: 49456
    Payload (36 bytes):
00000  ff fe 25 ff fb 03 ff fd 1f ff fd 21 ff fe 22 ff ..%......!... .
00016  fb 05 ff fa 20 01 ff f0 ff fa 27 01 ff f0 ff fa .... ..!.... .
00032  18 01 ff f0 .....

Packet number 10:
    From: 104.39.76.237
        To: 104.39.142.37
    Protocol: TCP
    Src port: 49456
    Dst port: 23

Capture complete.
william@ubuntu:~/lab4$
```

Sniffex will need some modification for us to sniff passwords. However, for short passwords like those password to the vm, there is no need for us to just output the payload.

Questions:

Question 4:

Yes since we can set the package length field to any value. But when we receive the package, it is always added to the zero that is not specified as packet size. Therefore, the payload is larger than the load and it will result in zero padding until the total length of the package is completed.

Question 5:

Yes since when we use the original socket programming in c, we must manually calculate the checksum, otherwise the error will be received. If we use the python scapy package, then it has a built-in function for us which is .checksum.

Question 6:

The original socket requires privileges that are not present in the normal user so it must be the root privilege user to run the command to start the original socket. It failed when it requested kernel access to create the socket.

## Task 2

### Task 2 (a)

We find that the packets that has been caught by tcpdump. We use a spoofing program and changed the ip address and results are shown below:

The screenshot shows two terminal windows side-by-side. The left window displays the output of the command `sudo tcpdump -v icmp`, capturing ICMP echo requests from 104.39.142.37 to google-public-dns-a.google.com. The right window shows the output of a spoofing program, which sent two ICMP echo requests to the same destination.

```
william@ubuntu:~/lab4/task2$ sudo tcpdump -v icmp
tcpdump: listening on ens33, link-type EN10MB (Ethernet), capture size 262144 bytes
07:29:15.589986 IP (tos 0x0, ttl 111, id 111, offset 0, flags [none], proto ICMP (1), length 128)
    104.39.142.37 > google-public-dns-a.google.com: ICMP echo request, id 50723, seq 26437, length 108
07:29:17.798493 IP (tos 0x0, ttl 111, id 111, offset 0, flags [none], proto ICMP (1), length 128)
    104.39.142.37 > google-public-dns-a.google.com: ICMP echo request, id 50723, seq 26437, length 108
]

william@ubuntu:~/lab4/task2$ sudo ./a.out ens33
Packet sent successfully
william@ubuntu:~/lab4/task2$ sudo ./a.out ens33
Packet sent successfully
william@ubuntu:~/lab4/task2$ 
```

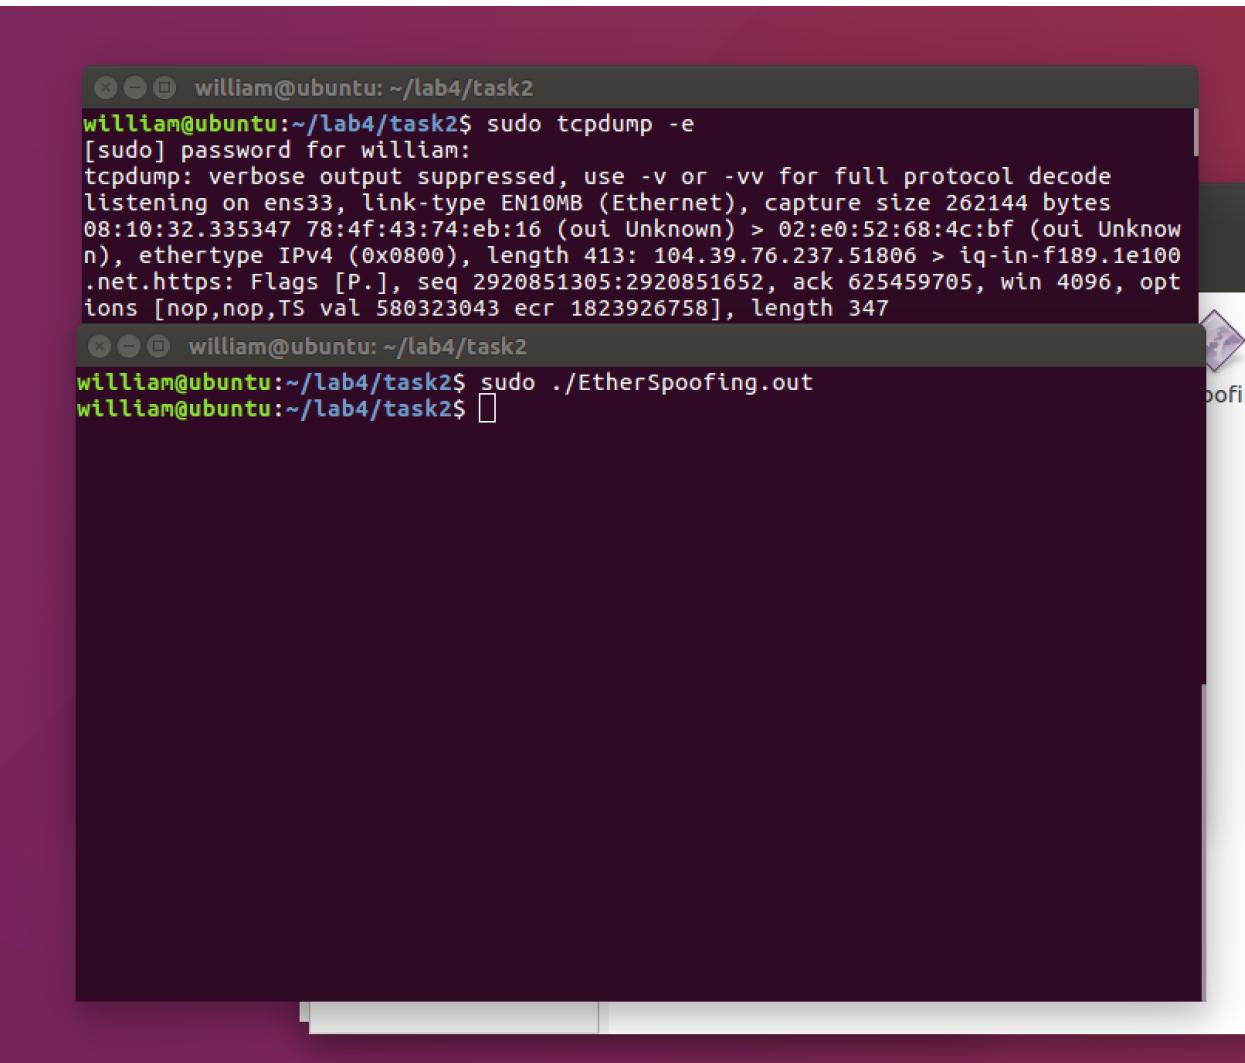
### Task2 (b)

The screenshot shows a terminal window displaying the output of the command `sudo tcpdump -v icmp`. It captures a large number of ICMP echo requests and replies between various IP addresses, indicating a Denial of Service (DoS) attack or a flood of traffic.

```
yanbo@yanbo-ThinkPad-Edge:~/443/lab4/task2$ sudo tcpdump -v icmp
tcpdump: listening on wlp3s0, link-type EN10MB (Ethernet), capture size 262144 bytes
11:01:17.969378 IP (tos 0x60, ttl 111, id 111, offset 0, flags [none], proto ICMP (1), length 128)
    104.39.142.37 > 104.39.120.152: ICMP echo request, id 50723, seq 26437, length 108
11:01:17.975274 IP (tos 0x60, ttl 64, id 19515, offset 0, flags [none], proto ICMP (1), length 128)
    104.39.120.152 > 104.39.142.37: ICMP echo reply, id 50723, seq 26437, length 108
11:01:37.790066 IP (tos 0x60, ttl 111, id 111, offset 0, flags [none], proto ICMP (1), length 128)
    104.39.142.37 > 104.39.120.152: ICMP echo request, id 50723, seq 26437, length 108
11:01:37.790117 IP (tos 0x60, ttl 64, id 21824, offset 0, flags [none], proto ICMP (1), length 128)
    104.39.120.152 > 104.39.142.37: ICMP echo reply, id 50723, seq 26437, length 108
11:01:39.732598 IP (tos 0x60, ttl 111, id 111, offset 0, flags [none], proto ICMP (1), length 128)
    104.39.142.37 > 104.39.120.152: ICMP echo request, id 50723, seq 26437, length 108
11:01:39.732645 IP (tos 0x60, ttl 64, id 21855, offset 0, flags [none], proto ICMP (1), length 128)
    104.39.120.152 > 104.39.142.37: ICMP echo reply, id 50723, seq 26437, length 108
11:01:40.887813 IP (tos 0x60, ttl 111, id 111, offset 0, flags [none], proto ICMP (1), length 128)
    104.39.142.37 > 104.39.120.152: ICMP echo request, id 50723, seq 26437, length 108
11:01:40.887858 IP (tos 0x60, ttl 64, id 22120, offset 0, flags [none], proto ICMP (1), length 128)
    104.39.120.152 > 104.39.142.37: ICMP echo reply, id 50723, seq 26437, length 108
```

We spoofed an ICMP echo request packet on behalf of another machine and the packet was sent to the ip address of 104.39.120.152. The tcpdump caught requested packets and thus replied packets.

### Task 2(c)



```
william@ubuntu:~/lab4/task2
william@ubuntu:~/lab4/task2$ sudo tcpdump -e
[sudo] password for william:
tcpdump: verbose output suppressed, use -v or -vv for full protocol decode
listening on ens33, link-type EN10MB (Ethernet), capture size 262144 bytes
08:10:32.335347 78:4f:43:74:eb:16 (oui Unknown) > 02:e0:52:68:4c:bf (oui Unknown), ethertype IPv4 (0x0800), length 413: 104.39.76.237.51806 > iq-in-f189.1e100.net.https: Flags [P.], seq 2920851305:2920851652, ack 625459705, win 4096, options [nop,nop,TS val 580323043 ecr 1823926758], length 347
william@ubuntu:~/lab4/task2
william@ubuntu:~/lab4/task2$ sudo ./EtherSpoofing.out
william@ubuntu:~/lab4/task2$
```

The source code is compiled and tested in the system. We successfully got the frames by tcpdump.

### Task 3

In this task, we use `inet addr()`, `inet network()`, `inet ntoa()`, `inet aton()` to convert IP addresses from the dotted decimal form (a string) to a 32-bit integer of network/host byte order.

The ip address of virtual machine A is 192.168.105.133. The ip address of virtual machine B is 192.168.105.132. Non-existing ip address is 192.168.105.134. When we ping a non-existing IP in the same LAN, an ARP request will be sent first. If there is no reply, the ICMP requests will not be sent later. So in order to avoid that ARP request which will stop the ICMP packet, we change the ARP cache of the victim VM by adding another MAC to the IP address mapping entry.

In this program, whenever it sees an ICMP echo request packet, it spoofs an ICMP echo reply packet. Therefore, even if the victim machine pings a non-existing machine, it will always see that the machine is alive.

In the program shown below, the ping's host and a sniffing program with spoofing facility. Basically my spoof program sniff an ICMP request packet on the network and make a ICMP reply packet and send it to the source of the ICMP request packet. When we ping the link, we are able to see the packets sent and received from VM A to VM B in wireshark and terminal, thus prove that our program can let VM B(192.168.105.132) relies VM A(192.168.105.133) by using non-existing host (192.168.105.134).

```
william@ubuntu: ~/lab4/task3
Sniffex - Sniffer example using libpcap
Copyright (c) 2005 The Tcpdump Group
THERE IS ABSOLUTELY NO WARRANTY FOR THIS PROGRAM.

Device: ens33
Number of packets: -1
Filter expression: icmp[icmptype]==icmp-echo

Packet number 1:
From: 192.168.105.133
To: 192.168.105.134
Protocol: ICMP
New Source address: 192.168.105.133
New Destination address: 192.168.105.133
==

packet sent
Sent 28 byte packet to 192.168.105.133

Packet number 2:
From: 192.168.105.133
To: 192.168.105.134
Protocol: ICMP
New Source address: 192.168.105.133
New Destination address: 192.168.105.133
```

Capturing from ens33

No.	Time	Source	Destination	Protocol	Length	Info
1	0.000000000	192.168.105.133	192.168.105.134	ICMP	98	Echo (ping) request
2	0.109833484	192.168.105.134	192.168.105.133	ICMP	60	Echo (ping) reply
3	1.001804371	192.168.105.133	192.168.105.134	ICMP	98	Echo (ping) request
4	1.134528543	192.168.105.134	192.168.105.133	ICMP	60	Echo (ping) reply
5	2.006847523	192.168.105.133	192.168.105.134	ICMP	98	Echo (ping) request
6	2.157890591	192.168.105.134	192.168.105.133	ICMP	60	Echo (ping) reply
7	3.030170602	192.168.105.133	192.168.105.134	ICMP	98	Echo (ping) request
8	3.182865882	192.168.105.134	192.168.105.133	ICMP	60	Echo (ping) reply
9	4.054612589	192.168.105.133	192.168.105.134	ICMP	98	Echo (ping) request
10	4.205800103	192.168.105.134	192.168.105.133	ICMP	60	Echo (ping) reply
11	5.079120560	192.168.105.133	192.168.105.134	ICMP	98	Echo (ping) request
12	5.230127022	Vmware_f9:cc:40	Vmware_c7:04:5b	ARP	60	Who has 192.168.105.
13	5.230155296	Vmware_c7:04:5b	Vmware_f9:cc:40	ARP	42	192.168.105.133 is a
14	5.231025180	192.168.105.134	192.168.105.133	ICMP	60	Echo (ping) reply
15	6.102280104	192.168.105.133	192.168.105.134	ICMP	98	Echo (ping) request
16	6.253408127	192.168.105.134	192.168.105.133	ICMP	60	Echo (ping) reply
17	7.126662898	192.168.105.133	192.168.105.134	ICMP	98	Echo (ping) request

```

0000 aa aa aa aa aa 00 0c 29 c7 04 5b 08 00 45 00 ..... )...[..E.
0010 00 54 da f6 40 00 40 01 0b 56 c0 a8 69 85 c0 a8 .T..@. @. .V..i...
0020 69 86 08 00 d7 07 26 10 00 01 0c b3 be 5a 00 00 i.....& .....Z..
0030 00 00 6f 06 02 00 00 00 00 00 10 11 12 13 14 15 ..o..... .
0040 16 17 18 19 1a 1b 1c 1d 1e 1f 20 21 22 23 24 25 ..... !%"$%

```

Our code as below:

```

#define APP_NAME "Sniffex"
#define APP_DESC "Sniffer example using libpcap"
#define APP_COPYRIGHT "Copyright (c) 2005 The Tcpdump Group"
#define APP_DISCLAIMER "THERE IS ABSOLUTELY NO WARRANTY FOR THIS
PROGRAM."
#include <pcap.h>
#include <stdio.h>
#include <string.h>
#include <stdlib.h>
#include <ctype.h>
#include <errno.h>
#include <sys/types.h>
#include <sys/socket.h>
#include <netinet/in.h>
#include <arpa/inet.h>
#include <netdb.h>
#include <linux/ip.h>
#include <linux/icmp.h>
#include <string.h>
#include <unistd.h>

unsigned short in_cksum(unsigned short *, int);


```

```

void parse_argvs(char**, char*, char* );
void usage();
char* getip();
char* toip(char*);
void sendicmp(char *src_addr,char * dst_addr)
{
    struct iphdr* ip;
    struct iphdr* ip_reply;
    struct icmphdr* icmp;
    struct sockaddr_in connection;
    char* packet;
    char* buffer;
    int sockfd;
    int one=1;
    int *optval=&one;
    int addrlen;
    int siz;
/*
 * allocate all necessary memory
*/
printf("New Source address: %s\n", dst_addr);
printf("New Destination address: %s\n", src_addr);
packet = malloc(sizeof(struct iphdr) + sizeof(struct icmphdr));
buffer = malloc(sizeof(struct iphdr) + sizeof(struct icmphdr));
/****************************************/
ip = (struct iphdr*) packet;
icmp = (struct icmphdr*) (packet + sizeof(struct iphdr));
/**
here the ip packet is set up
*/
ip->ihl = 5;
ip->version = 4;
ip->tos = 0;
ip->tot_len = sizeof(struct iphdr) + sizeof(struct icmphdr);
ip->id = htons(0);
ip->frag_off= 0;
ip->ttl= 64;
ip->protocol= IPPROTO_ICMP;
ip->saddr = inet_addr("10.20.103.90");
ip->daddr = inet_addr(src_addr);
ip->check =0;
ip->check = in_cksum((unsigned short *)ip, sizeof(struct iphdr));
if ((sockfd = socket(AF_INET, SOCK_RAW, IPPROTO_ICMP)) == -1)

```

```

    {
        perror("socket");
        exit(EXIT_FAILURE);
    }
/*
 * IP_HDRINCL must be set on the socket so that
 * the kernel does not attempt to automatically add
 * a default ip header to the packet
*/
// optval=1;
setsockopt(sockfd, IPPROTO_IP, IP_HDRINCL, &optval, sizeof(int));
/*
 * here the icmp packet is created* also the ip checksum is generated
*/
icmp->type = 0;
icmp->code = 0;
icmp->un.echo.id = random();
icmp->un.echo.sequence= 0;
icmp->checksum=0;
icmp-> checksum= in_cksum((unsigned short *)icmp, sizeof(struct icmphdr));
ip->check= in_cksum((unsigned short *)packet, sizeof(struct iphdr)+sizeof(struct
icmphdr));
connection.sin_family = AF_INET;
connection.sin_addr.s_addr = ip->daddr;
//connection.sin_addr._addr = inet_addr(src_addr);
printf("\npacket sent\n");
/*
now the packet is sent
*/
//memcpy(packet+20,&icmp,8);
if( (sendto(sockfd, packet, ip->tot_len, 0, (struct sockaddr *)&connection, sizeof(struct
sockaddr)))== -1)
{
    perror("socket");
    exit(EXIT_FAILURE);
}
printf("Sent %d byte packet to %s\n", ip->tot_len, src_addr);
/*
*
now we listen for responses
*/
free(packet);
free(buffer);

```

```

        close(sockfd);
//return 0;
}
void usage()
{
    fprintf(stderr, "\nUsage: pinger [destination] <-s [source]>\n");
    fprintf(stderr, "Destination must be provided\n");
    fprintf(stderr, "Source is optional\n\n");
}
/*
* return the ip address if host provided by DNS name
*/
char* toip(char* address)
{
    struct hostent* h;
    h = gethostbyname(address);
    return inet_ntoa(*(struct in_addr *)h->h_addr);
}
/*
* in_cksum --
* Checksum routine for Internet Protocol
* family headers (C Version)
*/
unsigned short in_cksum(unsigned short *addr, int len)
{
    register int sum = 0;
    u_short answer = 0;
    register u_short *w = addr;
    register int nleft = len;
    /** Our algorithm is simple, using a 32 bit accumulator (sum), we add
     * sequential 16 bit words to it, and at the end, fold back all the
     * carry bits from the top 16 bits into the lower 16 bits.
    */
    while (nleft > 1)
    {
        sum += *w++;
        nleft -= 2;
    }
    /* mop up an odd byte, if necessary */
    if (nleft == 1)
    {
        *(u_char *) (&answer) = *(u_char *) w;
        sum += answer;
    }
}

```

```

/* add back carry outs from top 16 bits to low 16 bits */
sum = (sum >> 16) + (sum & 0xffff); /* add hi 16 to low 16 */
sum += (sum >> 16); /* add carry */
answer = ~sum; /* truncate to 16 bits */
    return (answer);
}
/* default snap length (maximum bytes per packet to capture) */
#define SNAP_LEN 1518
/* ethernet headers are always exactly 14 bytes [1] */
#define SIZE_ETHERNET 14
/* Ethernet addresses are 6 bytes */
#define ETHER_ADDR_LEN 6
/* Ethernet header */
struct sniff_ethernet {
    u_char ether_dhost[ETHER_ADDR_LEN]; /* destination host address */
    u_char ether_shost[ETHER_ADDR_LEN]; /* source host address */
    u_short ether_type;
/* IP? ARP? RARP?etc */
};

/* IP header */
struct sniff_ip {
    u_char ip_vhl;
/* version << 4 | header length >> 2 */
    u_char ip_tos;
/* type of service */
    u_short ip_len;
/* total length */
    u_short ip_id;
/* identification */
    u_short ip_off;
/* fragment offset field */
#define IP_RF 0x8000
/* reserved fragment flag */
#define IP_DF 0x4000
/* dont fragment flag */
#define IP_MF 0x2000
/* more fragments flag */
#define IP_OFFMASK 0x1fff
/* mask for fragmenting bits */
    u_char ip_ttl;
/* time to live */
    u_char ip_p;
/* protocol */
    u_short ip_sum;

```

```

/* checksum */
struct in_addr ip_src,ip_dst; /* source and dest address */
};

#define IP_HL(ip) ((ip->ip_vhl) & 0x0f)
#define IP_V(ip) ((ip->ip_vhl) >> 4)
/* TCP header */
typedef u_int tcp_seq;
struct sniff_tcp
{
    u_short th_sport; /* source port */
    u_short th_dport; /* destination port */
    tcp_seq th_seq; /* sequence number */
    tcp_seq th_ack; /* acknowledgement number */
    u_char th_offx2; /* data offset, rsvd */
#define TH_OFF(th) (((th)->th_offx2 & 0xf0) >> 4)
    u_char th_flags;
#define TH_FIN 0x01
#define TH_SYN 0x02
#define TH_RST 0x04
#define TH_PUSH 0x08
#define TH_ACK 0x10
#define TH_URG 0x20
#define TH_ECE 0x40
#define TH_CWR 0x80
#define TH_FLAGS (TH_FIN|TH_SYN|TH_RST|TH_ACK|TH_URG|TH_ECE|TH_CWR)
    u_short th_win;
/* window */
    u_short th_sum;
/* checksum */
    u_short th_urp;
/* urgent pointer */
};

//void got_packet(u_char *args, const struct pcap_pkthdr *header, const u_char *packet);
//void print_payload(const u_char *payload, int len);
//void print_hex_ascii_line(const u_char *payload, int len, int offset);
//void print_app_banner(void);
//void print_app_usage(void);
/*
 * app name/banner
 */
void print_app_banner(void)
{
    printf("%s - %s\n", APP_NAME, APP_DESC);
    printf("%s\n", APP_COPYRIGHT);
}

```

```

printf("%s\n", APP_DISCLAIMER);
printf("\n");
return;
}
/*
* print help text
*/
void print_app_usage(void)
{
    printf("Usage: %s [interface]\n", APP_NAME);
    printf("\n");
    printf("Options:\n");
    printf(" interface Listen on <interface> for packets.\n");
    printf("\n");
    return;
}
/*
* print data in rows of 16 bytes: offset hex ascii
*
* 00000 47 45 54 20 2f 20 48 54 54 50 2f 31 2e 31 0d 0a GET / HTTP/1.1..
*/
void print_hex_ascii_line(const u_char *payload, int len, int offset)
{
    int i;
    int gap;
    const u_char *ch;
/* offset */
    printf("%05d ", offset);
/* hex */
    ch = payload;
    for(i = 0; i < len; i++) {
        printf("%02x ", *ch);
        ch++;
    }
/* print extra space after 8th byte for visual aid */
    if (i == 7)
        printf(" ");
}
/* print space to handle line less than 8 bytes */
if (len < 8)
    printf(" ");
/* fill hex gap with spaces if not full line */
if (len < 16) {
    gap = 16 - len;
    for (i = 0; i < gap; i++) {

```

```

                printf(" ");
            }
        }
        printf(" ");
    /* ascii (if printable) */
    ch = payload;
    for(i = 0; i < len; i++) {
        if (isprint(*ch))
            printf("%c", *ch);
        else
            printf(".");
        ch++;
    }
    printf("\n");
    return;
}
/*
 * print packet payload data (avoid printing binary data)
*/
void print_payload(const u_char *payload, int len)
{
    int len_rem = len;
    int line_width = 16;
    int line_len;
    int offset = 0;
/* number of bytes per line */
/* zero-based offset counter */
    const u_char *ch = payload;
    if (len <= 0)
        return;
/* data fits on one line */
    if (len <= line_width) {
        print_hex_ascii_line(ch, len, offset);
        return;
    }
/* data spans multiple lines */
    for (;;) {
/* compute current line length */
        line_len = line_width % len_rem;
/* print line */
        print_hex_ascii_line(ch, line_len, offset);
/* compute total remaining */
        len_rem = len_rem - line_len;
/* shift pointer to remaining bytes to print */

```

```

        ch = ch + line_len;
/* add offset */
        offset = offset + line_width;
/* check if we have line width chars or less */
        if (len_rem <= line_width) {
/* print last line and get out */
        print_hex_ascii_line(ch, len_rem, offset);
        break;
    }
}
return;
}
*/
/* dissect/print packet
*/
void got_packet(u_char *args, const struct pcap_pkthdr *header, const u_char *packet)
{
//sendicmp/inet_ntoa(ip->ip_src),inet_ntoa(ip->ip_dst));
    static int count = 1;
/* packet counter */
/* declare pointers to packet headers */
const struct sniff_ether *ether; /* The ethernet header [1] */
    const struct sniff_ip *ip;
/* The IP header */
    const struct sniff_tcp *tcp;
/* The TCP header */
    const char *payload;
/* Packet payload */
    int size_ip;
    int size_tcp;
    int size_payload;
//sendicmp/inet_ntoa(ip->ip_src),inet_ntoa(ip->ip_dst));
    printf("\nPacket number %d:\n", count);
    count++;
/* define ethernet header */
    ether = (struct sniff_ether*)(packet);
/* define/compute ip header offset */
    ip = (struct sniff_ip*)(packet + SIZE_ETHERNET);
    size_ip = IP_HL(ip)*4;
    if (size_ip < 20)
    {
        printf(" * Invalid IP header length: %u bytes\n", size_ip);
        return;
    }
}

```

```

/* print source and destination IP addresses */
printf("From: %s\n", inet_ntoa(ip->ip_src));
printf("To: %s\n", inet_ntoa(ip->ip_dst));
/* determine protocol */
switch(ip->ip_p) {
    case IPPROTO_TCP:
        printf(" Protocol: TCP\n");
        break;
    case IPPROTO_UDP:
        printf(" Protocol: UDP\n");
        return;
    case IPPROTO_ICMP:
        printf(" Protocol: ICMP\n");
        sendicmp(inet_ntoa(ip->ip_src),inet_ntoa(ip->ip_dst));
        return;
    case IPPROTO_IP:
        printf(" Protocol: IP\n");
        return;
    default:
        printf(" Protocol: unknown\n");
        return;
}
/*
 * OK, this packet is TCP.
*/
/* define/compute tcp header offset */
tcp = (struct sniff_tcp*)(packet + SIZE_ETHERNET + size_ip);
size_tcp = TH_OFF(tcp)*4;
if (size_tcp < 20) {
    printf(" * Invalid TCP header length: %u bytes\n", size_tcp);
    return;
}
printf(" Src port: %d\n", ntohs(tcp->th_sport));
printf(" Dst port: %d\n", ntohs(tcp->th_dport));
/* define/compute tcp payload (segment) offset */
payload = (u_char*)(packet + SIZE_ETHERNET + size_ip + size_tcp);
/* compute tcp payload (segment) size */
size_payload = ntohs(ip->ip_len) - (size_ip + size_tcp);
/*
 * Print payload data; it might be binary, so don't just
 * treat it as a string.
*/
if (size_payload > 0) {
    printf(" Payload (%d bytes):\n", size_payload);
}

```

```

        print_payload(payload, size_payload);
    }
//sendicmp/inet_ntoa(ip->ip_src),inet_ntoa(ip->ip_dst));
    return;
}
int main(int argc, char **argv)
{
char *dev = NULL; /* capture device name */
char errbuf[PCAP_ERRBUF_SIZE]; /* error buffer */
pcap_t *handle; /* packet capture handle */
char *filter_exp=argv[1]; /* filter expression [3] */
struct bpf_program fp; /* compiled filter program (expression) */
bpf_u_int32 mask; /* subnet mask */
bpf_u_int32 net; /* ip */
int num_packets = 10; /* number of packets to capture */
    print_app_banner();
/* check for capture device name on command-line
if (argc == 2) {
dev = argv[1];
}
else if (argc > 2) {
fprintf(stderr, "error: unrecognized command-line options\n\n");
print_app_usage();
exit(EXIT_FAILURE);
}
else {
find a capture device if not specified on command-line */
    dev = pcap_lookupdev(errbuf);
    if (dev == NULL) {
        fprintf(stderr, "Couldn't find default device: %s\n",
               errbuf);
        exit(EXIT_FAILURE);
    }
/* get network number and mask associated with capture device */
    if (pcap_lookupnet(dev, &net, &mask, errbuf) == -1) {
        fprintf(stderr, "Couldn't get netmask for device %s: %s\n",
               dev, errbuf);
        net = 0;
        mask = 0;
    }
/* print capture info */
    printf("Device: %s\n", dev);
    printf("Number of packets: %d\n", num_packets);
    printf("Filter expression: %s\n", filter_exp);

```

```
/* open capture device */
handle = pcap_open_live(dev, SNAP_LEN, 1, 1000, errbuf);
if (handle == NULL) {
    fprintf(stderr, "Couldn't open device %s: %s\n", dev, errbuf);
    exit(EXIT_FAILURE);
}
/* make sure we're capturing on an Ethernet device [2] */
if (pcap_datalink(handle) != DLT_EN10MB) {
    fprintf(stderr, "%s is not an Ethernet\n", dev);
    exit(EXIT_FAILURE);
}
/* compile the filter expression */
if (pcap_compile(handle, &fp, filter_exp, 0, net) == -1) {
    fprintf(stderr, "Couldn't parse filter %s: %s\n",
            filter_exp, pcap_geterr(handle));
    exit(EXIT_FAILURE);
}
/* apply the compiled filter */
if (pcap_setfilter(handle, &fp) == -1) {
    fprintf(stderr, "Couldn't install filter %s: %s\n",
            filter_exp, pcap_geterr(handle));
    exit(EXIT_FAILURE);
}
/* now we can set our callback function */
pcap_loop(handle, num_packets, got_packet, NULL);
/* cleanup */
pcap_freecode(&fp);
pcap_close(handle);
printf("\nCapture complete.\n");
return 0;
}
```